

Introduction to AI

Lecture 6: Informed search strategies

Mona Taghavi



LaSalle College
Montréal

Informed vs uninformed search



- An informed search is a searching technique **that has additional information about the distance from the current state** to the goal.
- An informed search provides the AI with guidance on how and where to look for the solution to the problem. On the other hand, an uninformed search only provides the AI with the problem's specifications and no additional information about the potential solution.

Informed search



Informed search:

Use heuristics to guide the search

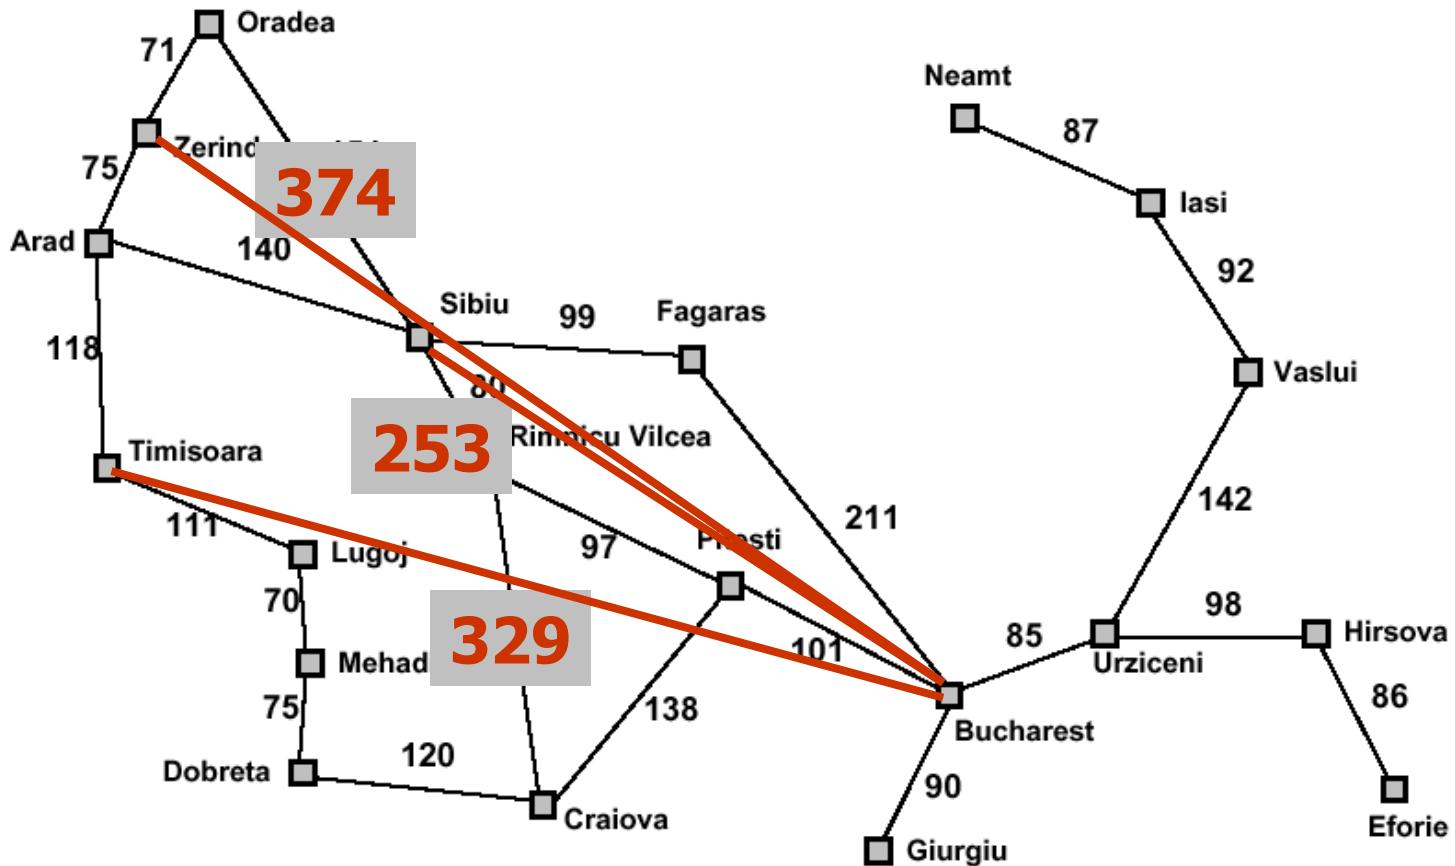
- A*
- Heuristics
- Hill-climbing
- Simulated annealing
- Genetic Algorithm
- Beam
- Etc.

Best-first search



- Idea:
 - use an evaluation function for each node; estimate of **"desirability"**
 - ⇒ expand most desirable unexpanded node.
- Implementation:
 - QueueingFn** = insert successors in decreasing order of desirability
- Special cases:
 - greedy search
 - A* search

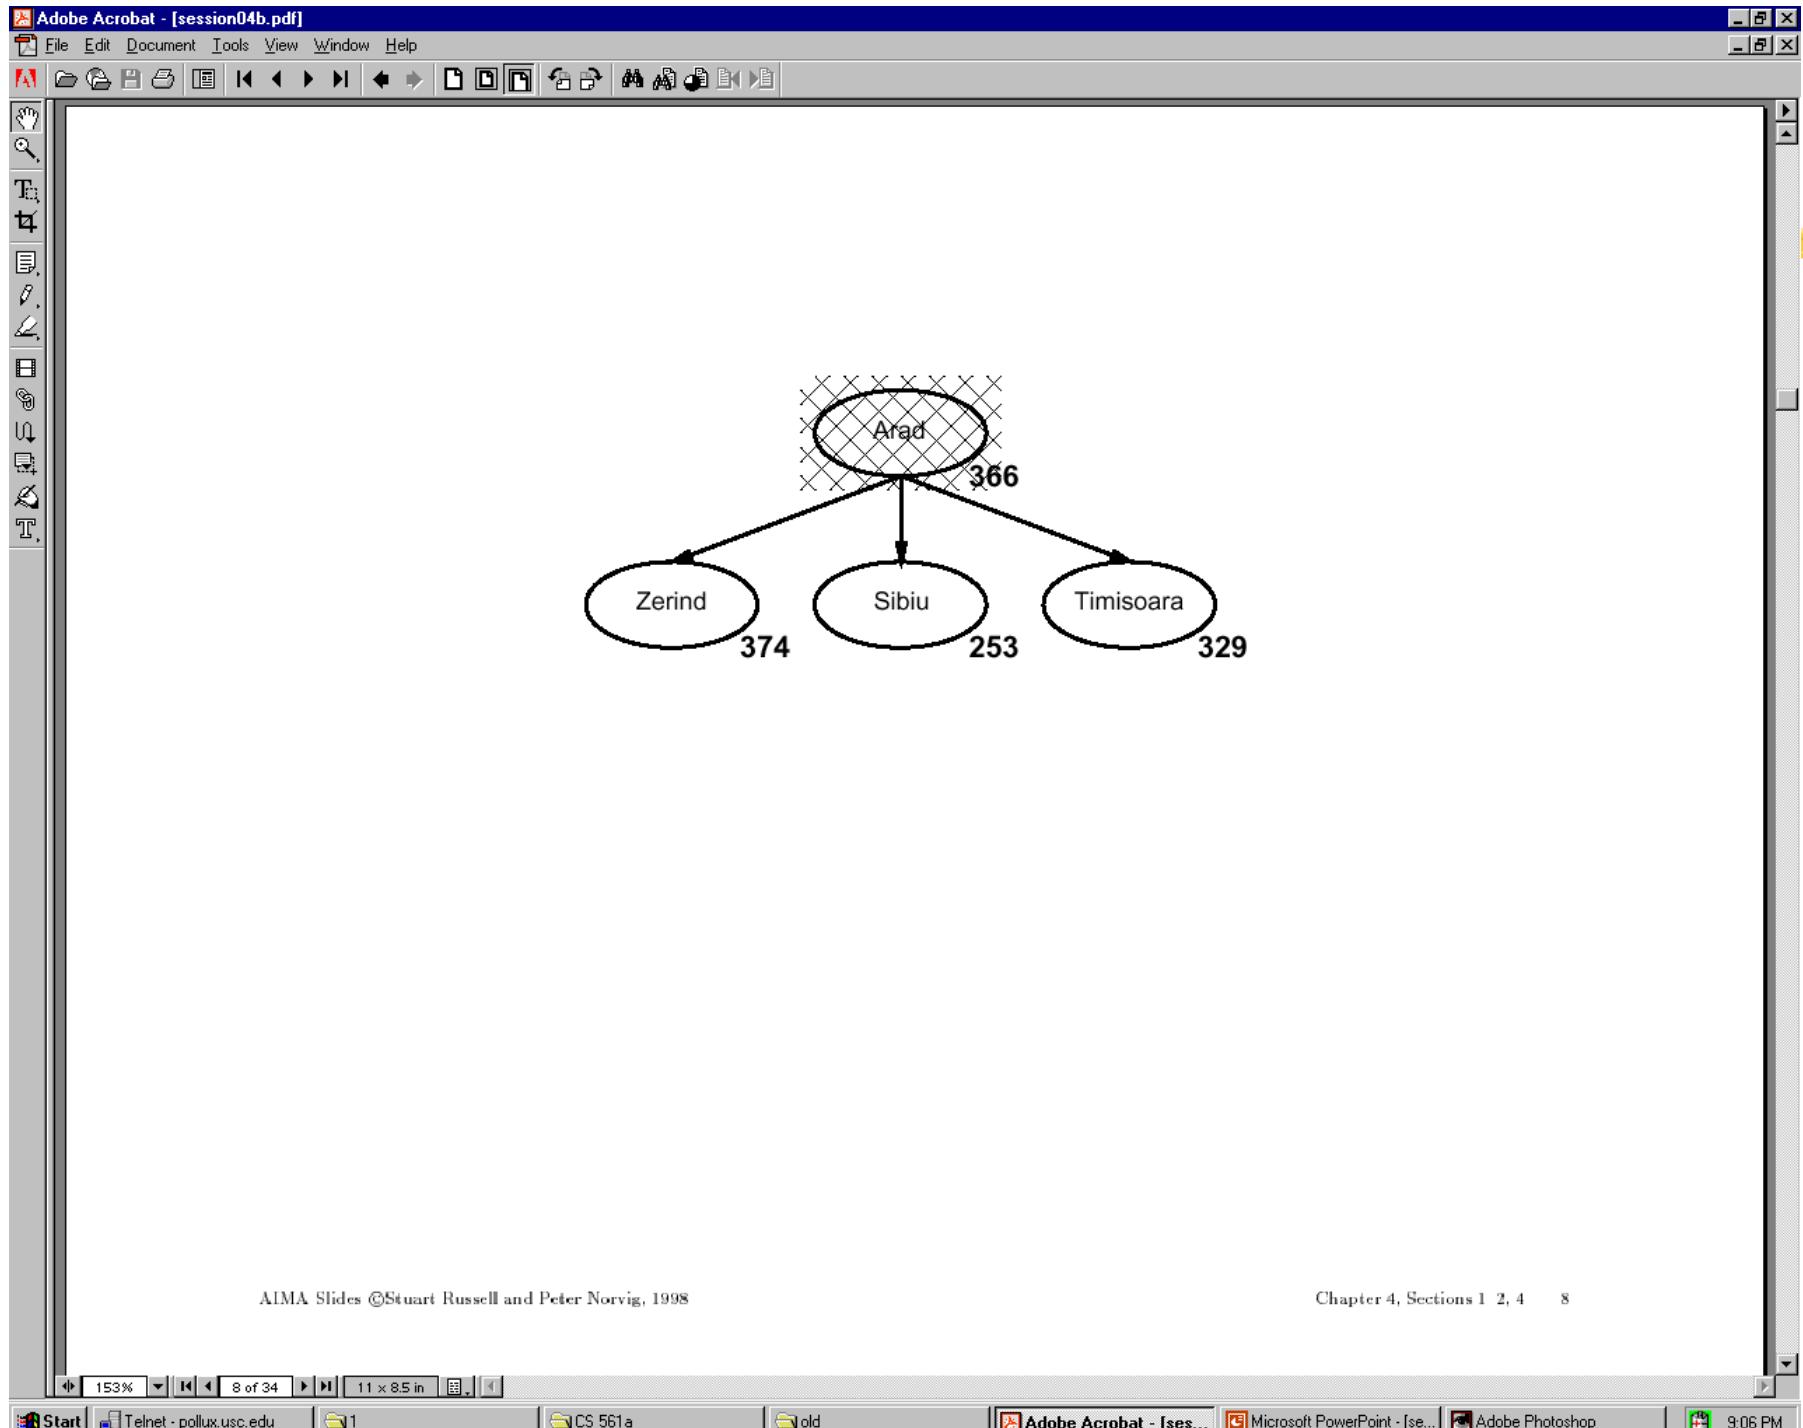
Romania with step costs in km

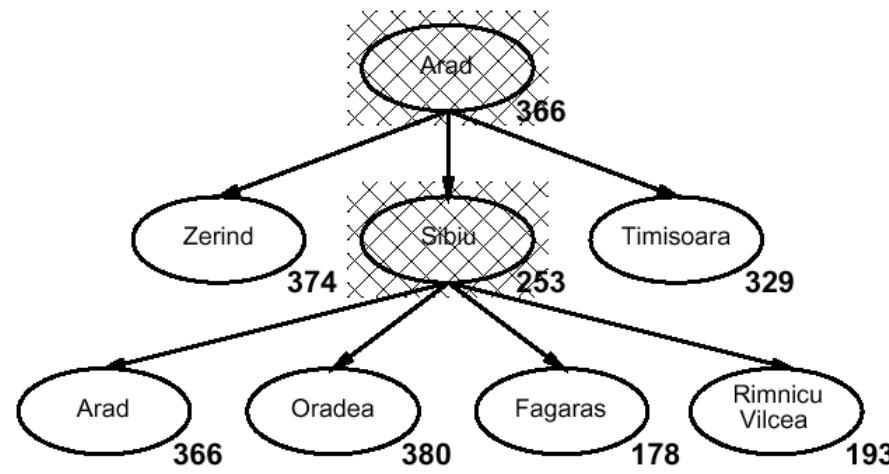


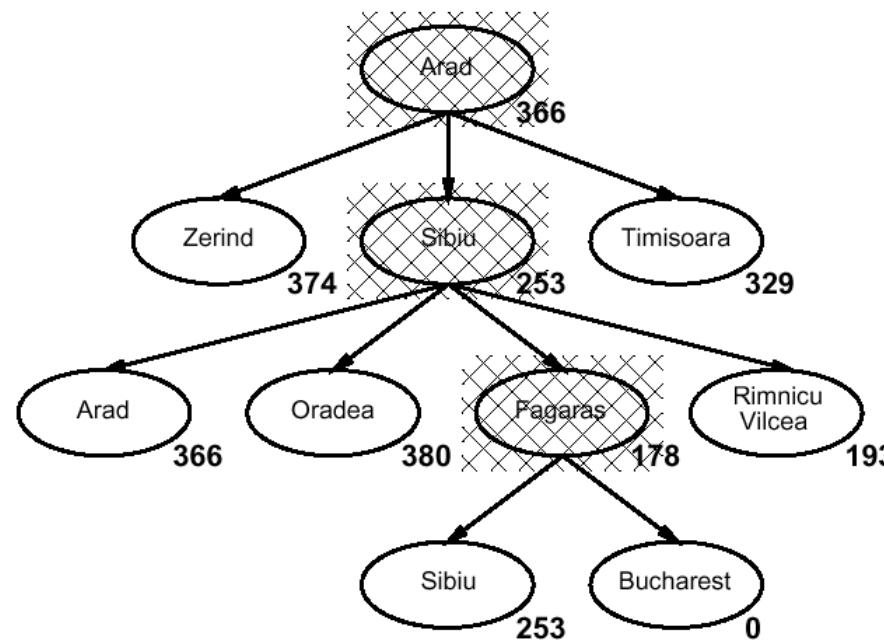
Greedy search

- Estimation function:
 $h(n)$ = estimate of distance from n to goal(heuristic)
- For example:
 $h_{SLD}(n)$ = straight-line distance from n to Bucharest
- Greedy search expands first the node that appears to be closest to the goal, according to $h(n)$.









Measuring Distance



Euclidean Distance

Manhattan Distance

Chebyshev Distance

A* search

- Idea: avoid expanding paths that are already expensive

evaluation function: $f(n) = g(n) + h(n)$ with:

$g(n)$ – cost so far to reach n

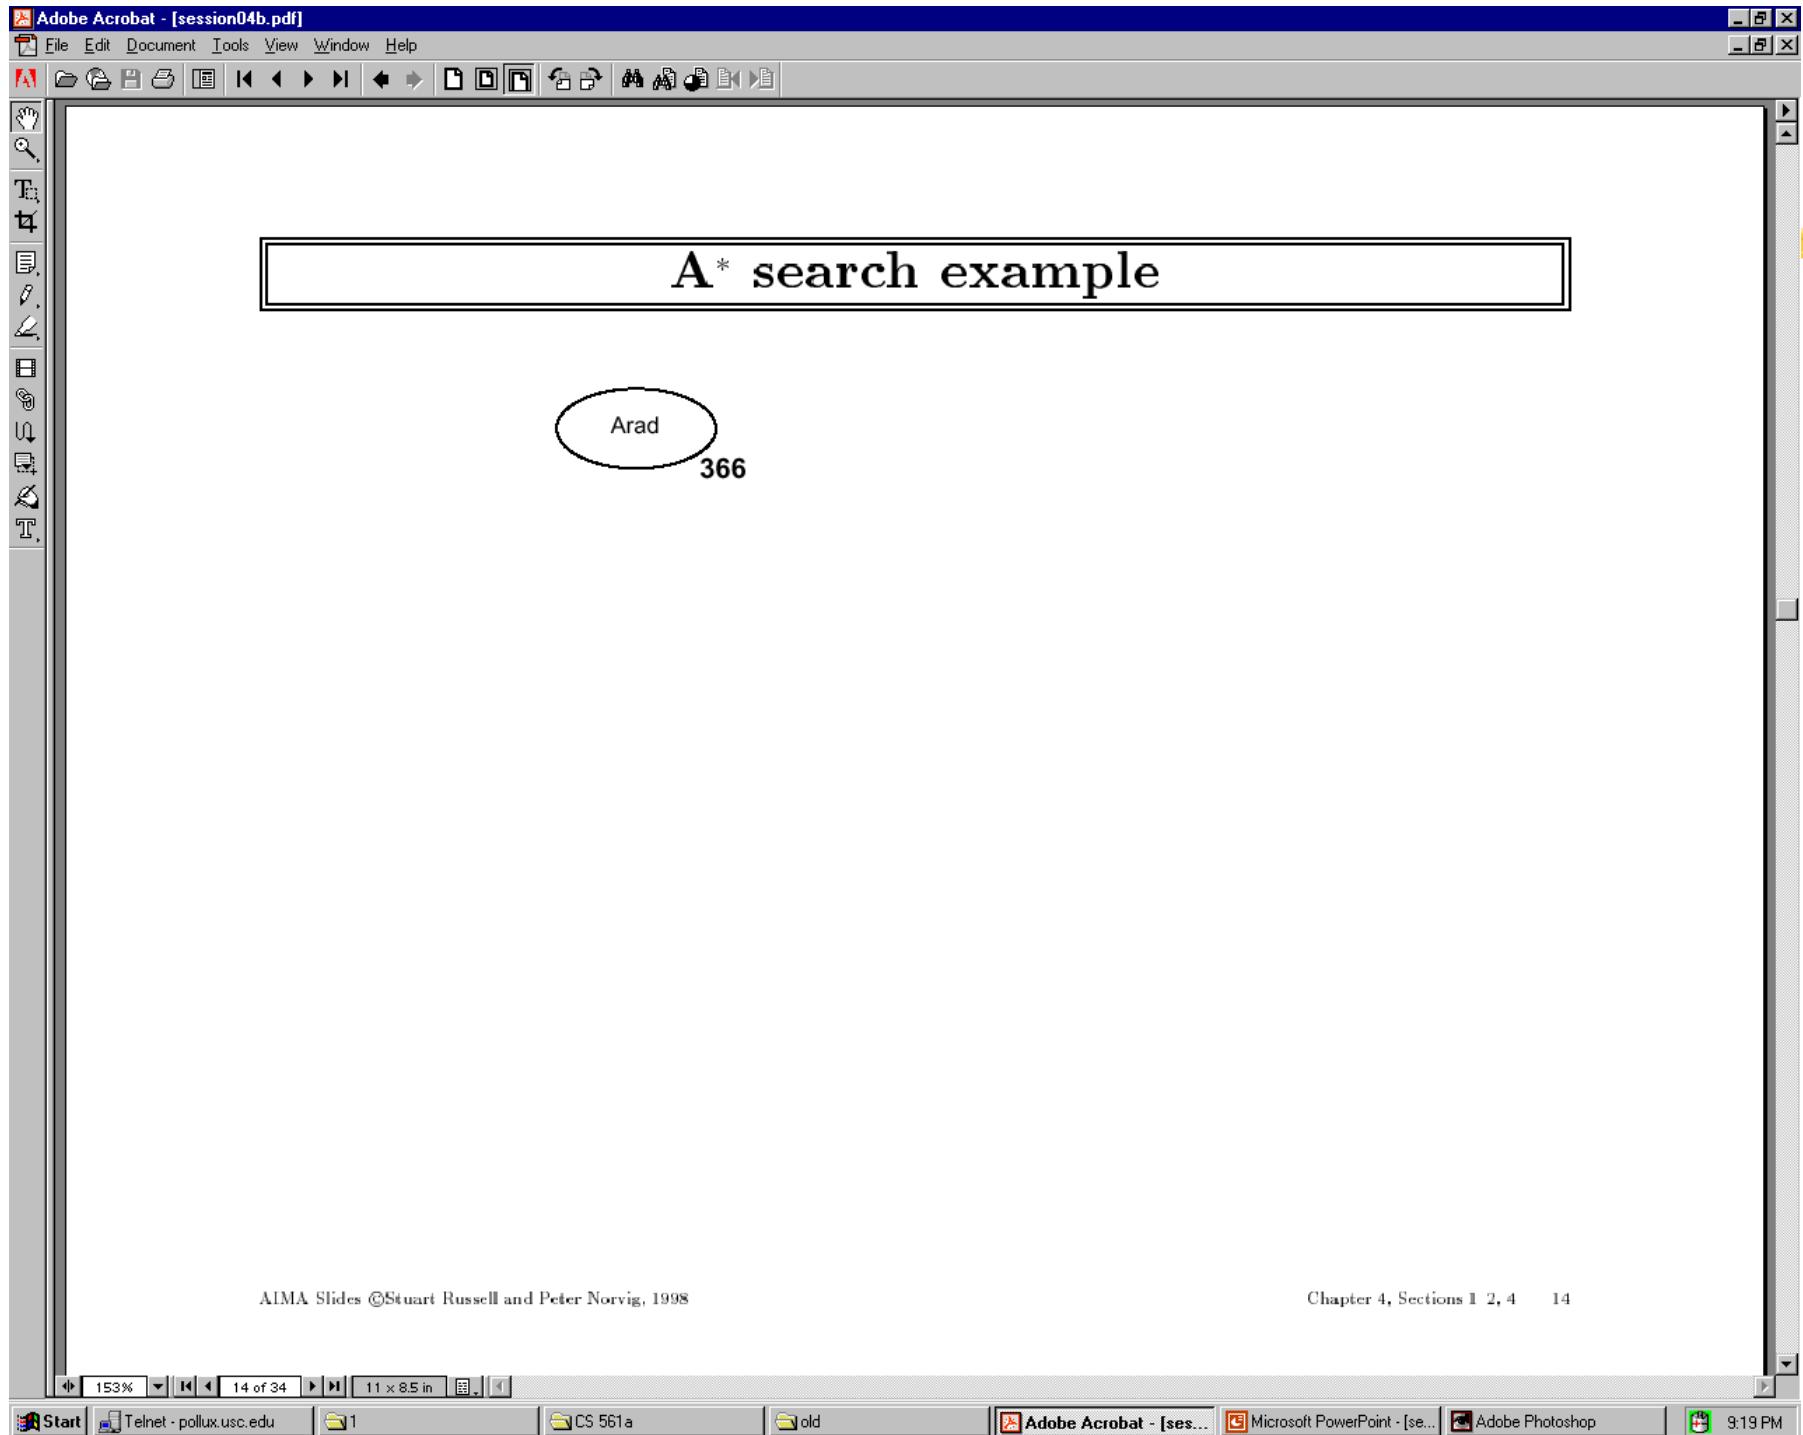
$h(n)$ – estimated cost to goal from n

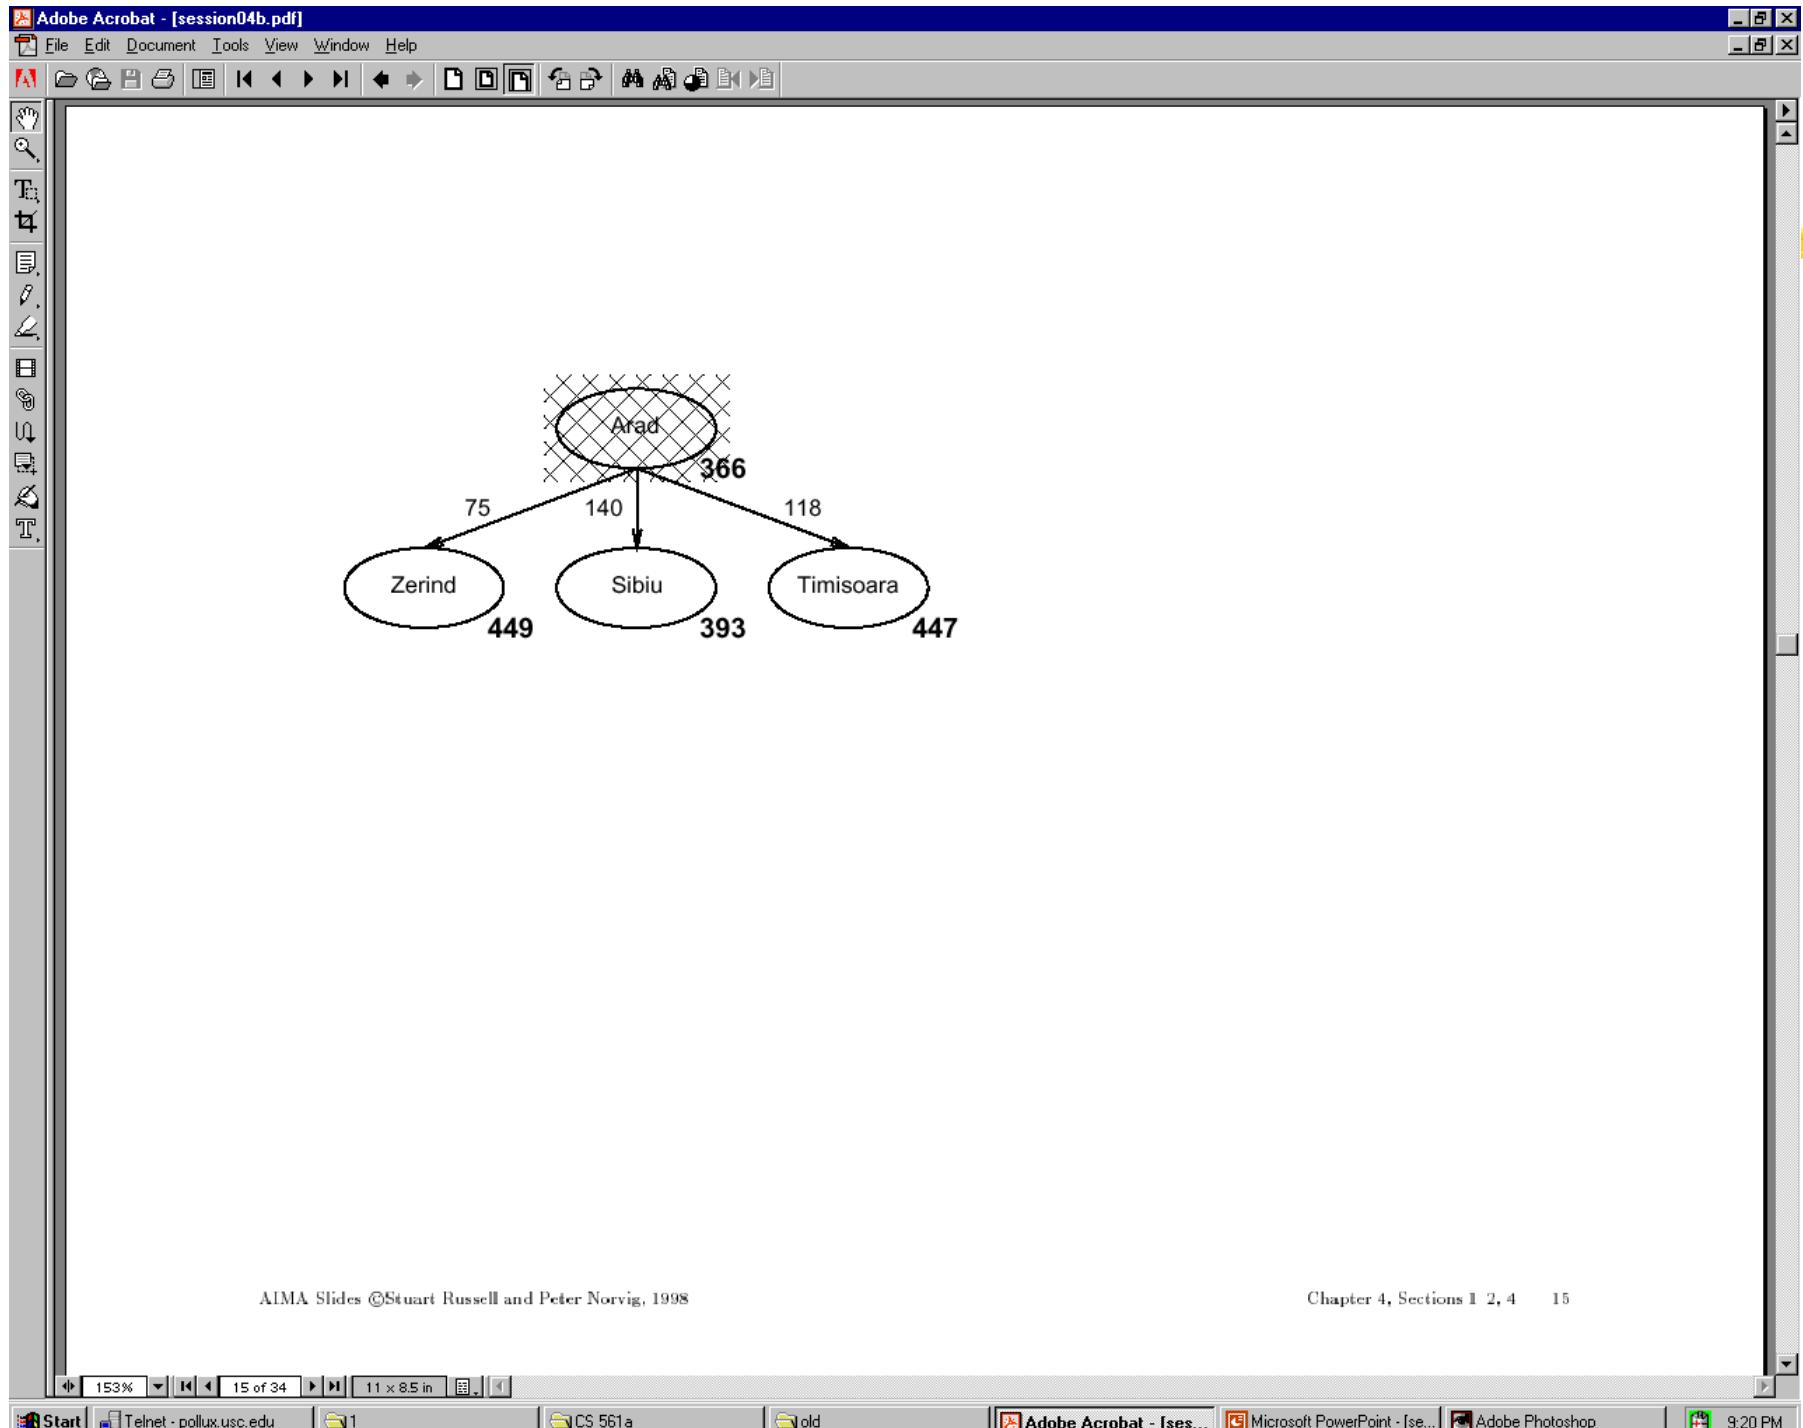
$f(n)$ – estimated total cost of path through n to goal

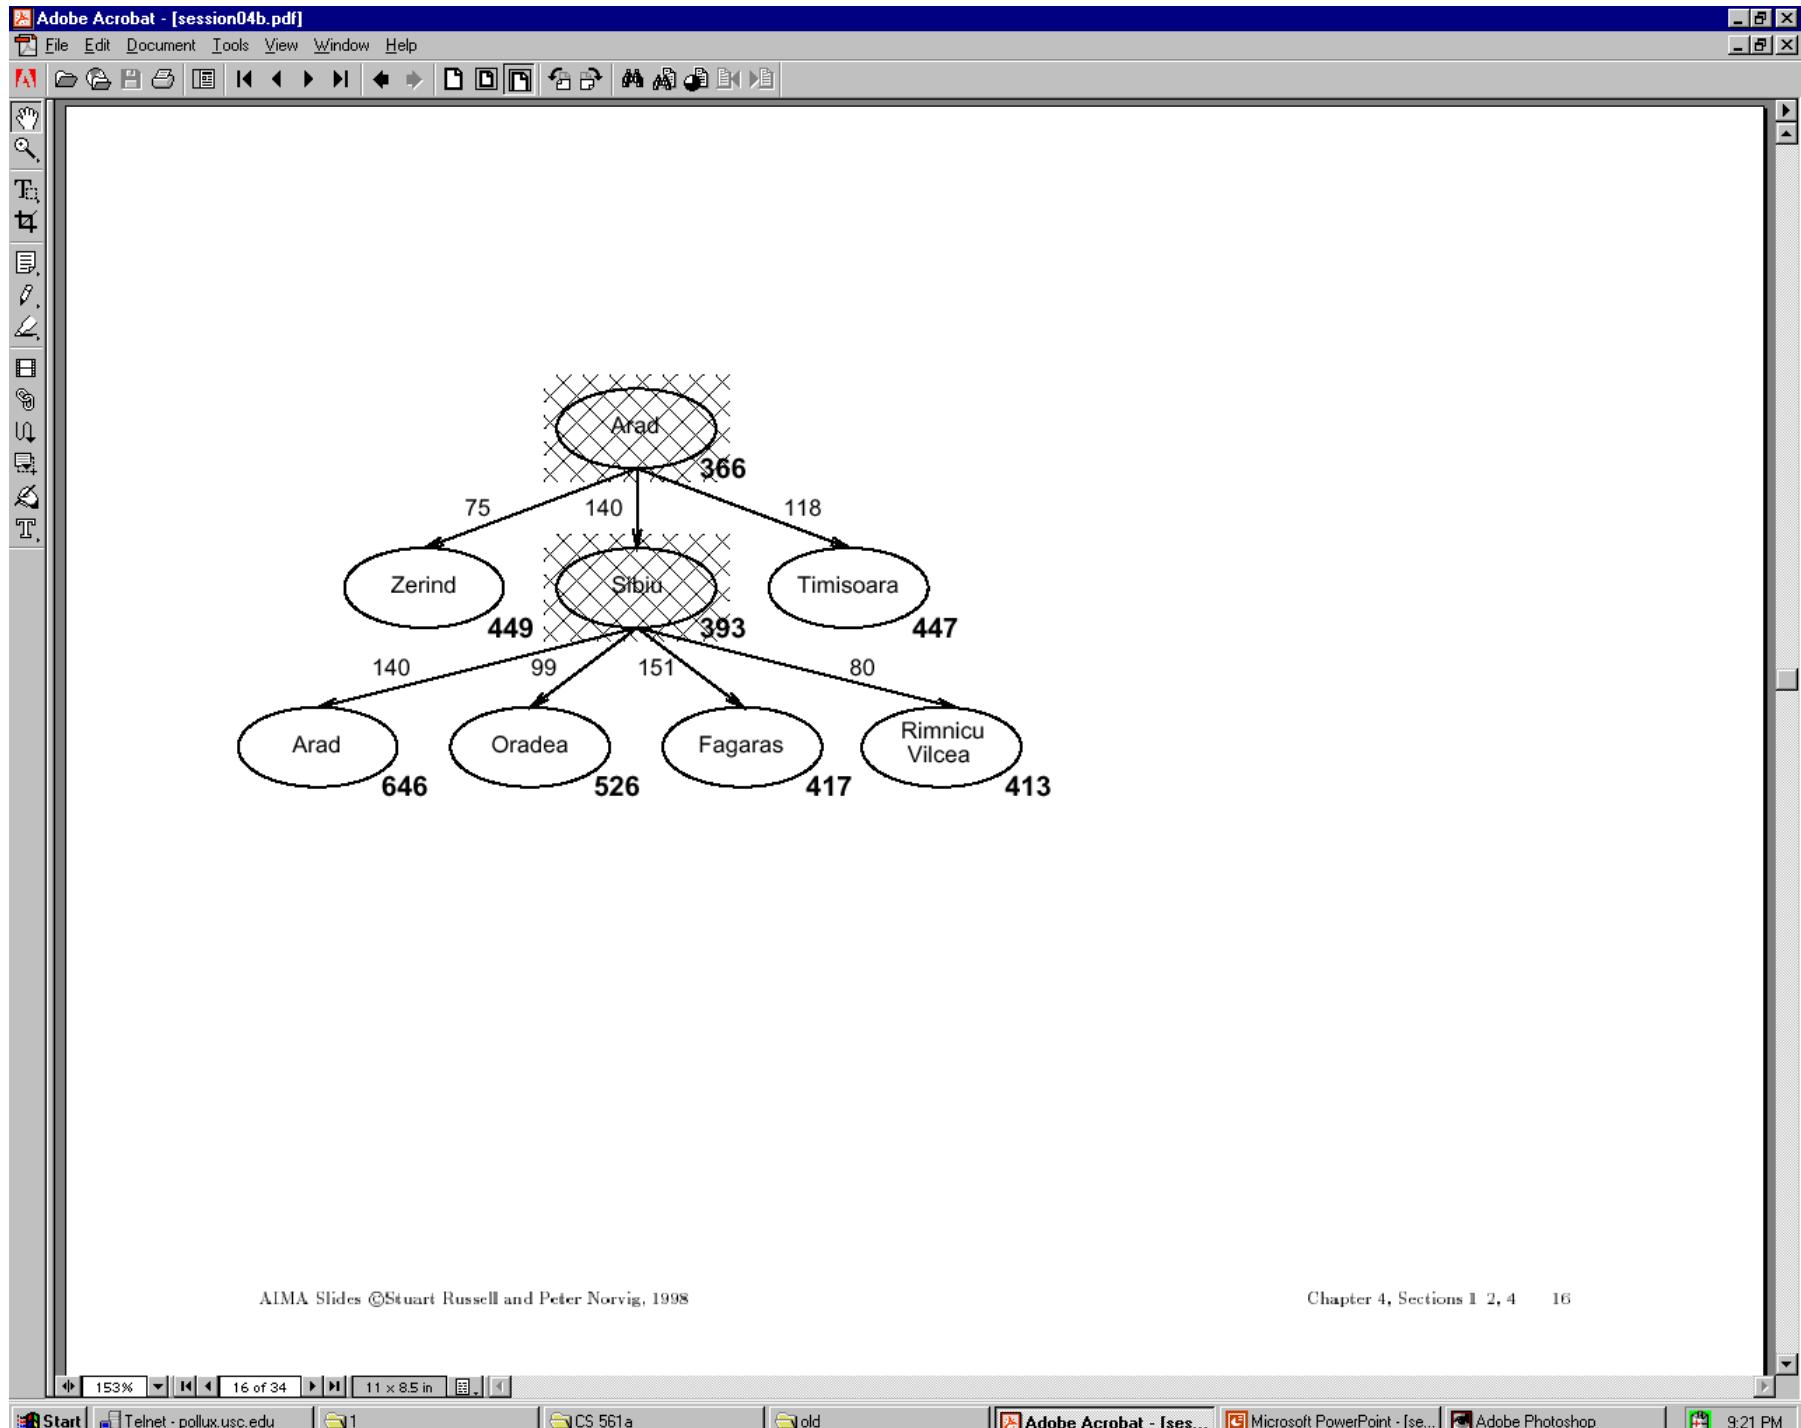
- A* search uses an admissible heuristic, that is,

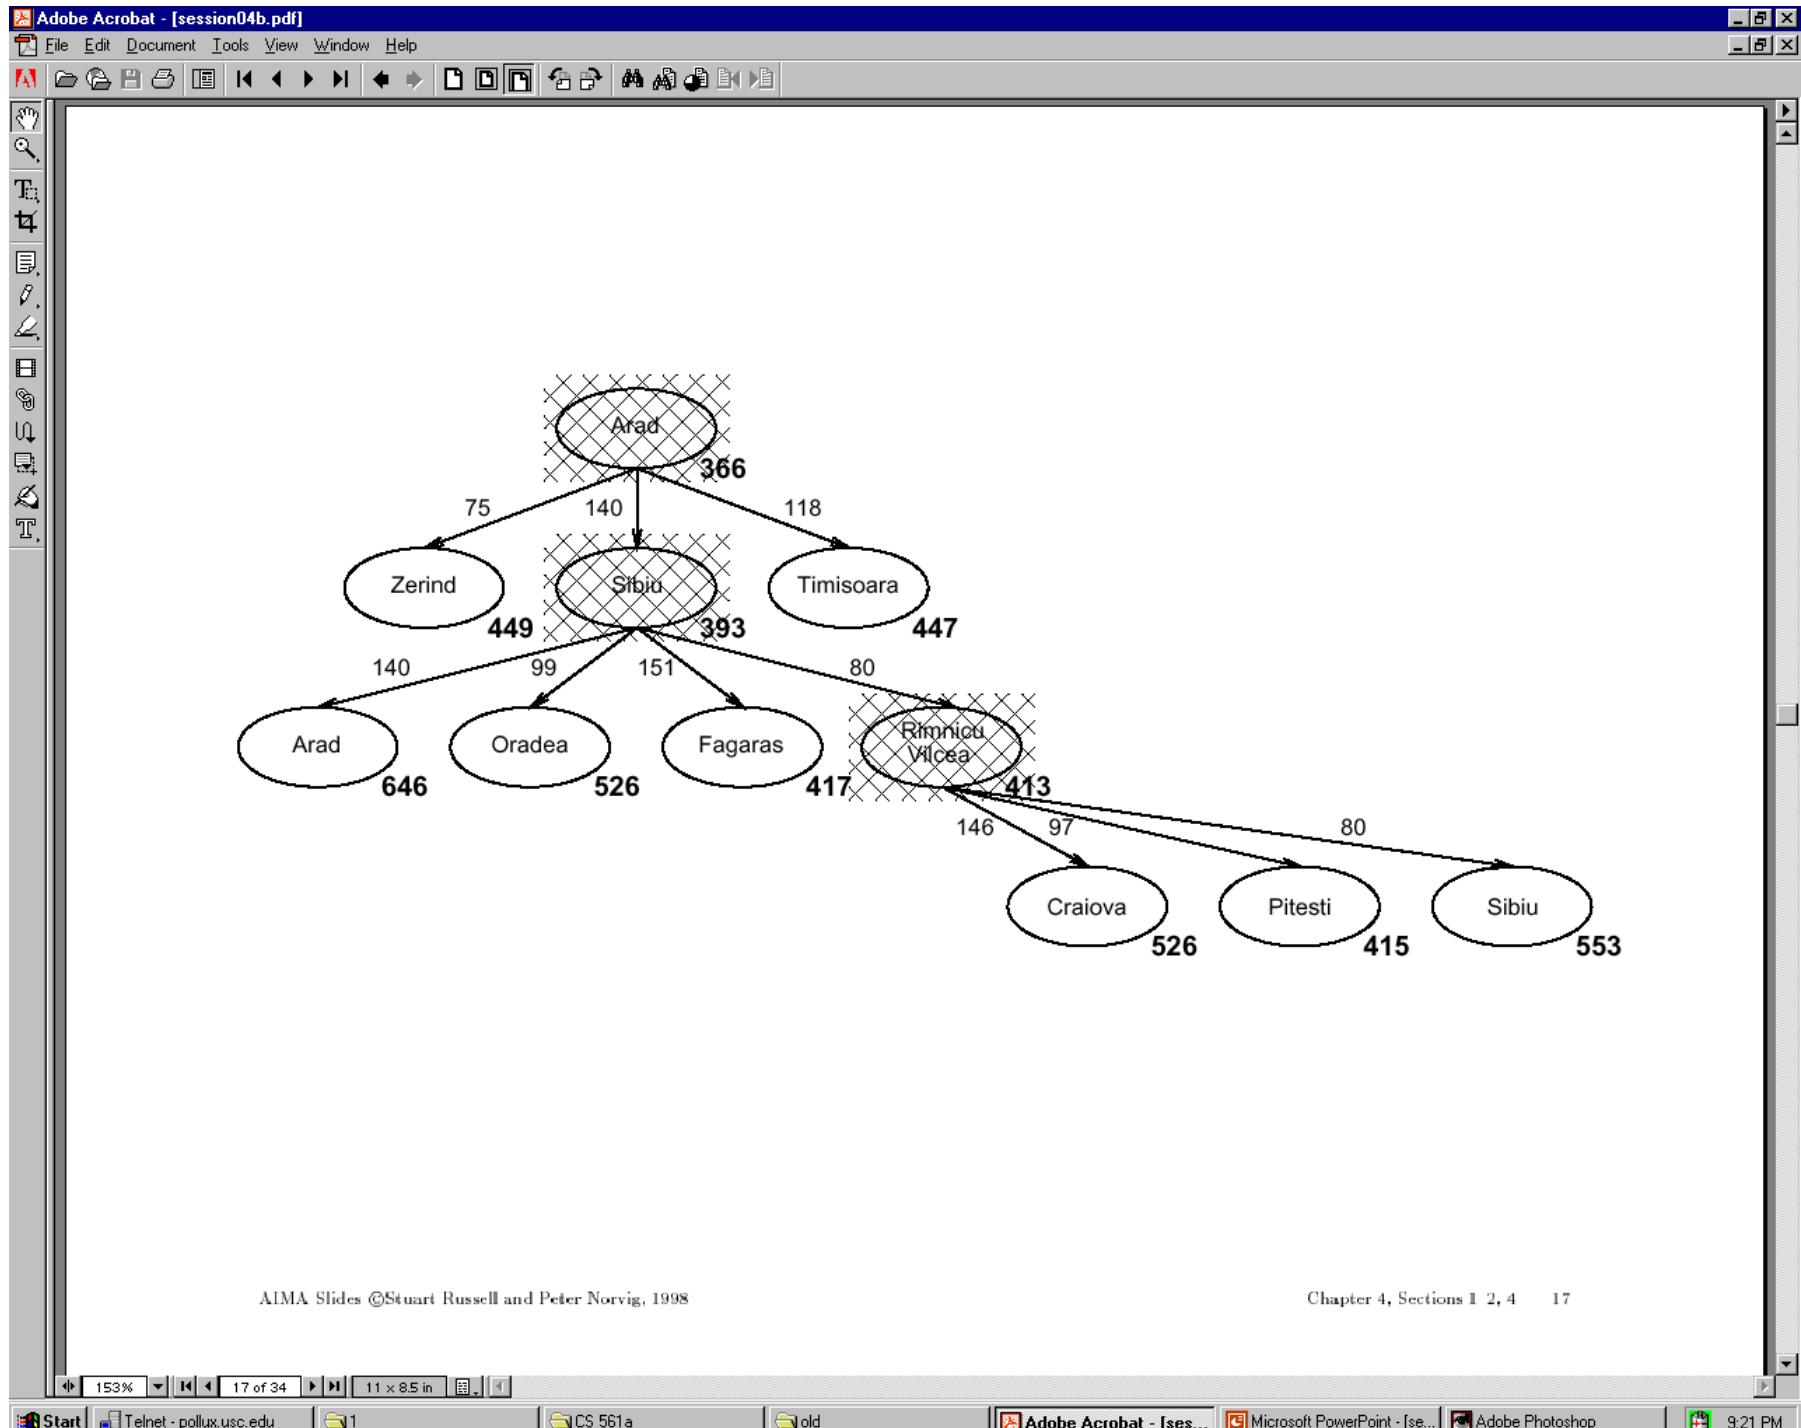
$h(n) \leq h^*(n)$ where $h^*(n)$ is the true cost from n .

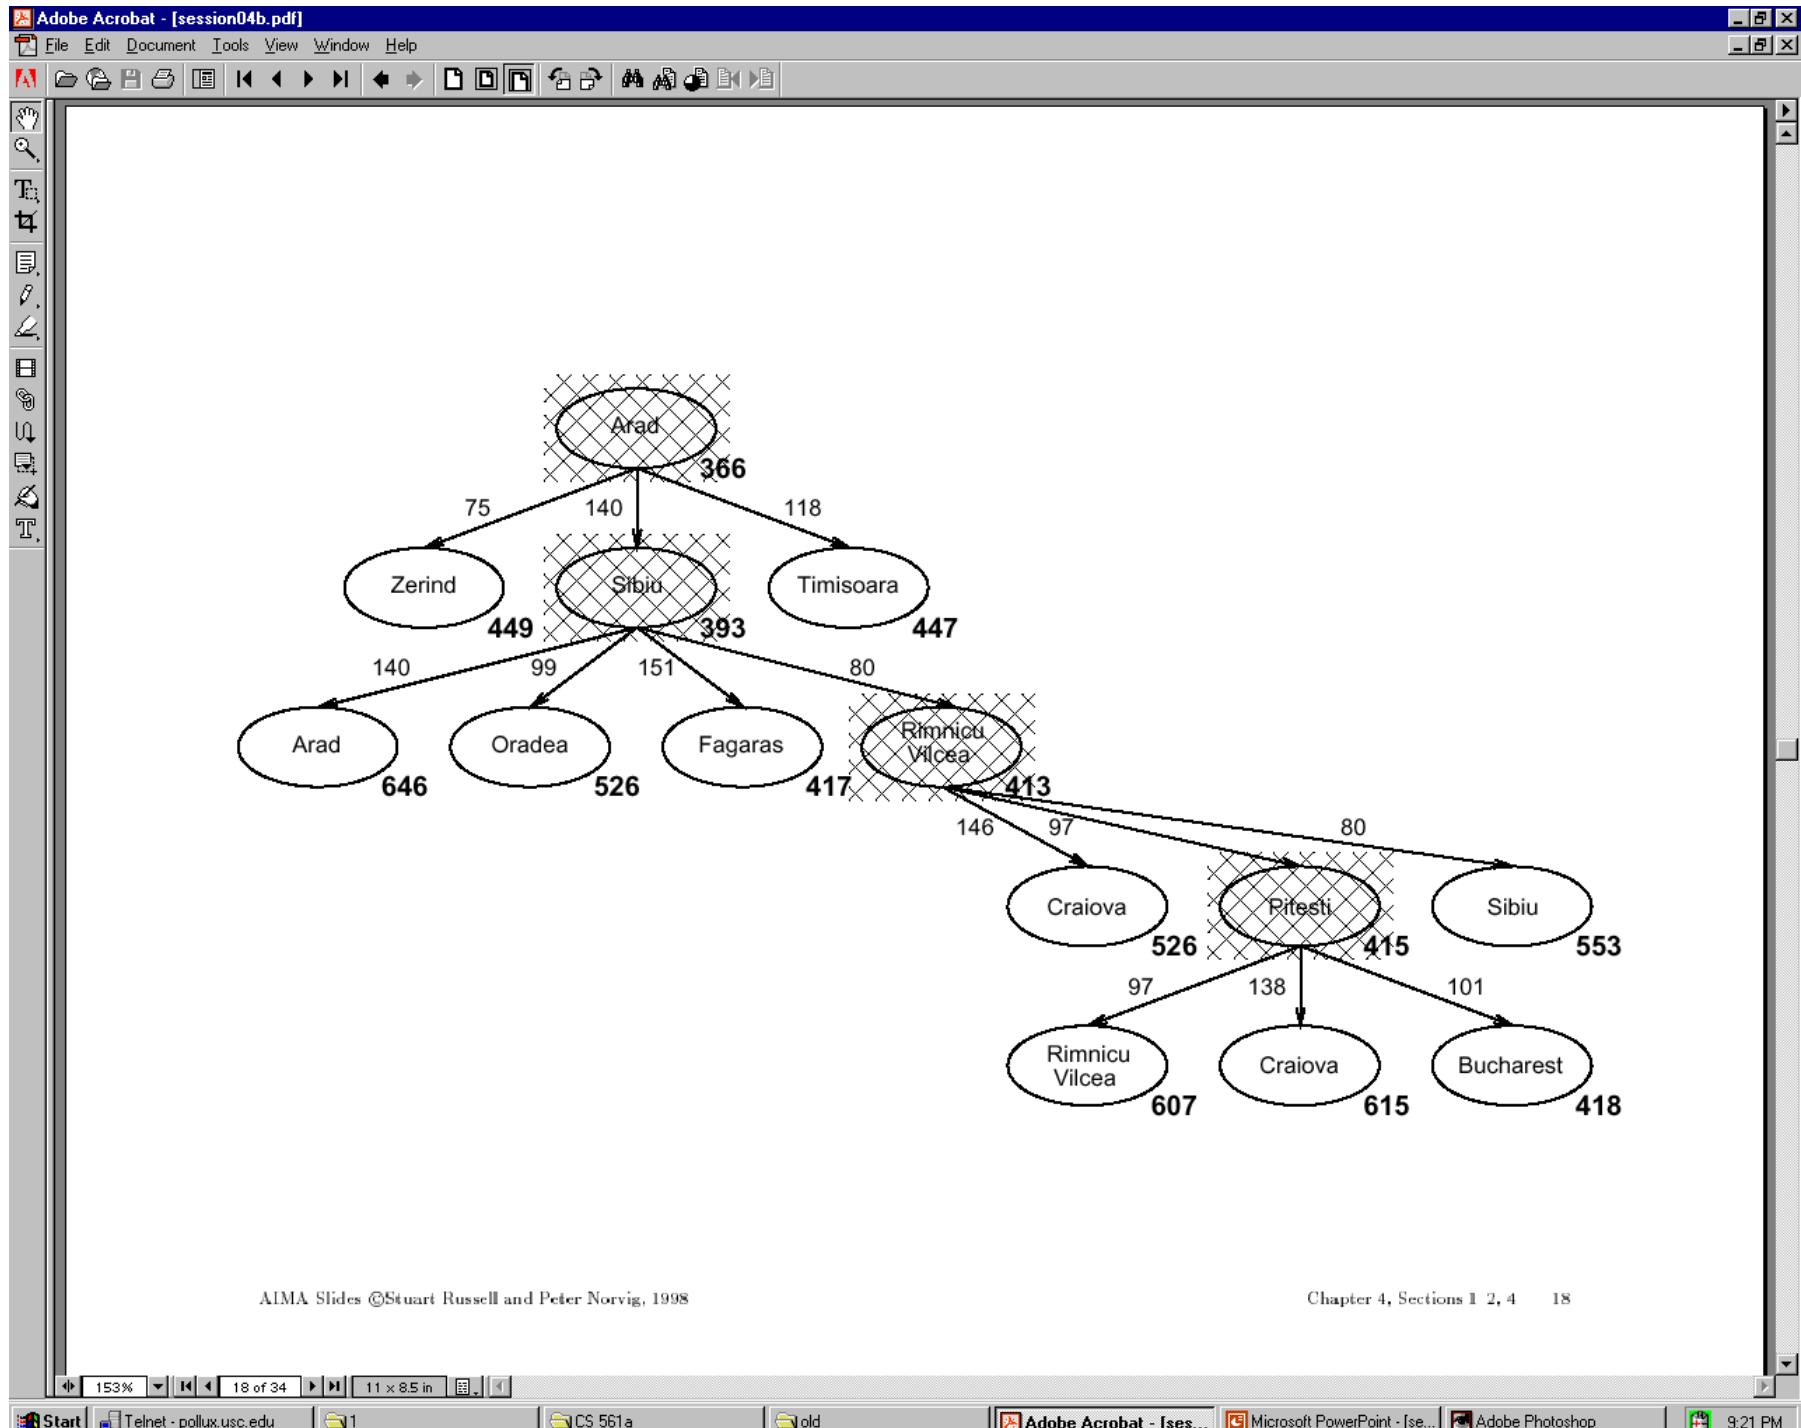
For example: $h_{SLD}(n)$ never overestimates actual road distance.

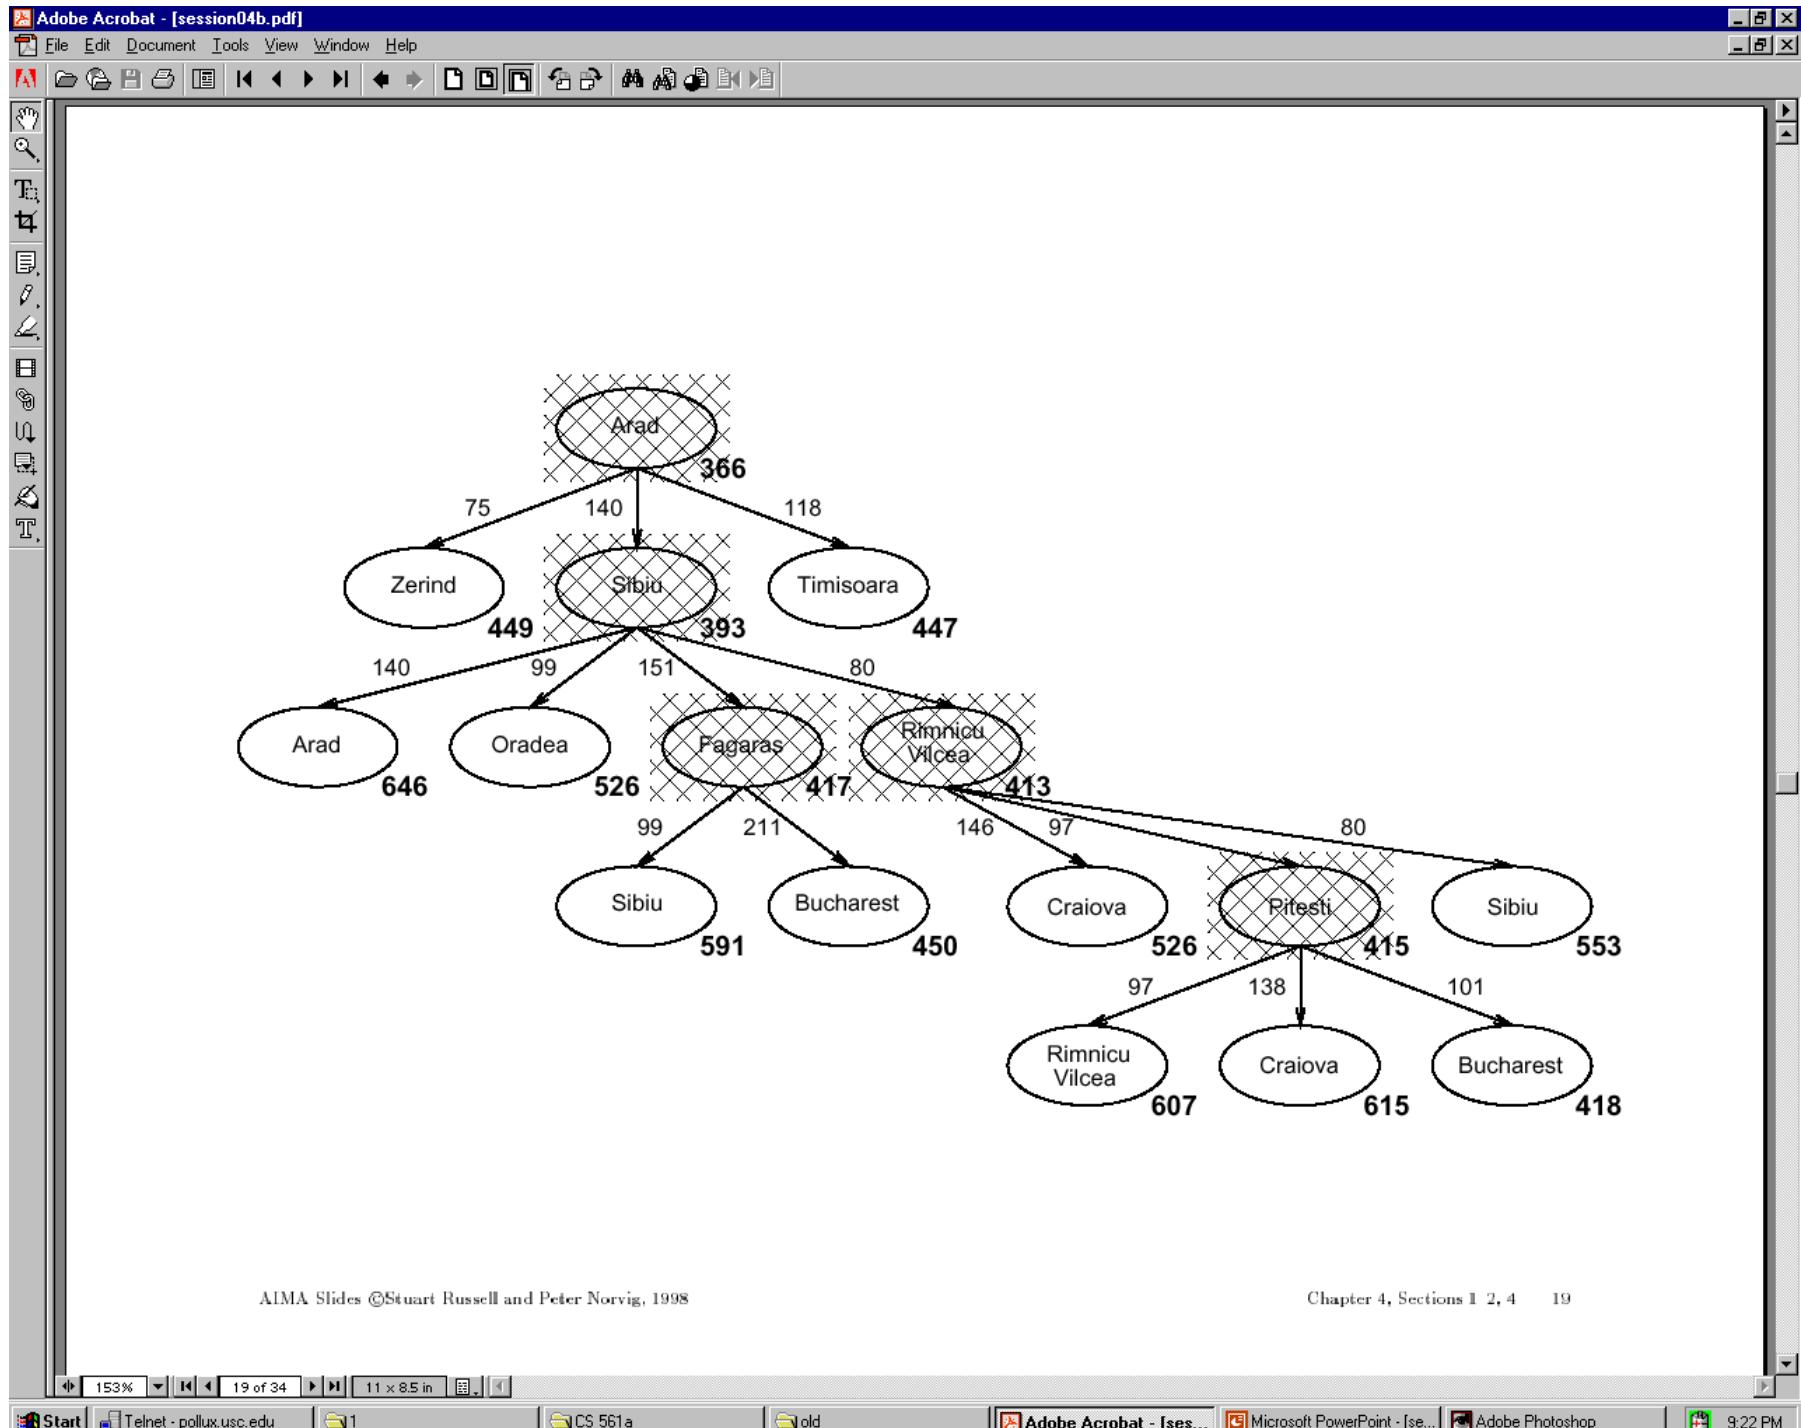












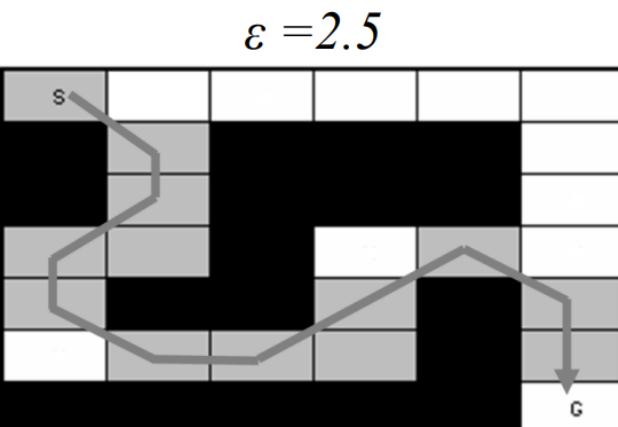
Weighted A* Search



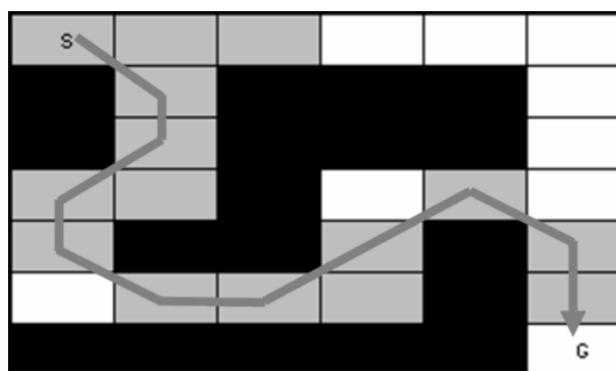
- Expands states in the order of $f(n) = g(n) + \epsilon h(n)$ values
 - $\epsilon > 1$ = bias towards states that are closer to goal
- Trades off optimality for speed
- In games we often multiply the heuristic by some constant factor to make A* run faster.

Effect of the Heuristic Function

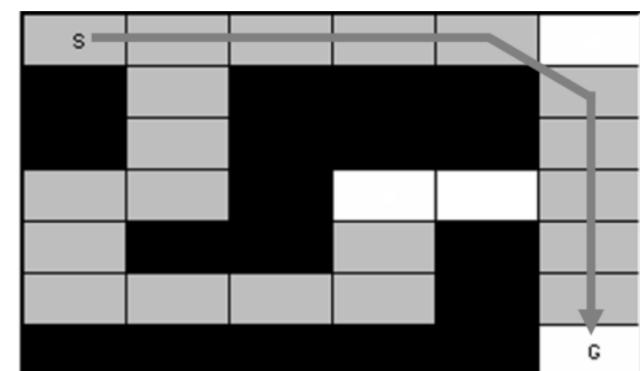
- Constructing anytime search based on weighted A*:
 - find the best path possible given some amount of time for planning
 - do it by running a series of weighted A* searches with decreasing ε



$\varepsilon = 2.5$
13 expansions
solution=11 moves



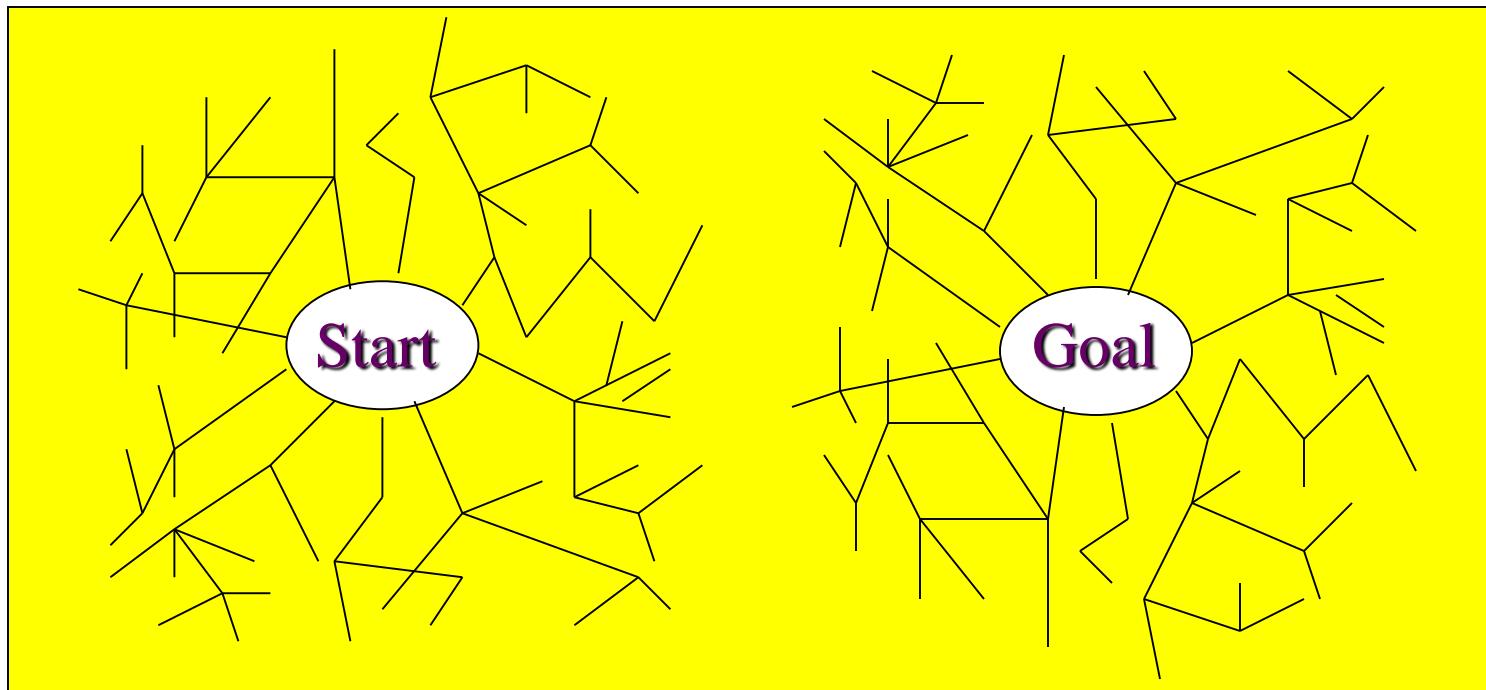
$\varepsilon = 1.5$
15 expansions
solution=11 moves



$\varepsilon = 1.0$
20 expansions
solution=10 moves

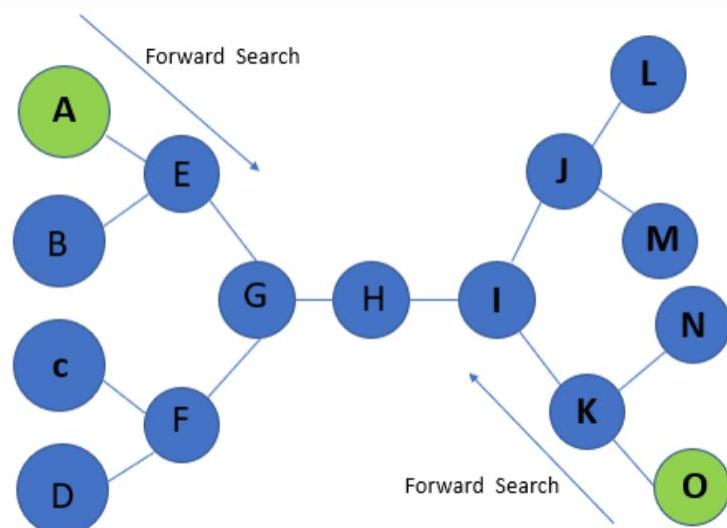
Bi-directional search (uninformed/informed search)

- Compute the tree both from the start-node and from a goal node, until these meet.



Bi-directional search

- Instead of searching from the start to the finish, you can start two searches in parallel—one from start to finish, and one from finish to start. When they meet, you should have a good path.
- The main idea behind bidirectional searches is to reduce the time taken for search drastically.
- It's a good idea that will help in some situations. The idea behind bidirectional searches is that searching results in a “tree” that fans out over the map. A big tree is much worse than two small trees, so it's better to have two small search trees.



Use BFS in Bidirectional search to move from A to H

